

Logic and Modelling

Todd Davies

January 15, 2015

Introduction

This is a unique course developed at the University of Manchester. It explains how implementations of logic can be used to solve a number a number of problems, such as solving hardest Sudoku puzzles in no time, analysing two-player games, or finding serious errors in computer systems.

Aims

This course intends to build an understanding of fundamentals of (mathematical) logic as well as some of the applications of logic in modern computer science, including hardware verification, finite domain constraint satisfaction and verification of concurrent systems.

Additional reading

Contents

- 1 Setting the stage
 - 1.1 Decision problems
 - 1.1.1 Decidable problems
 - 1.2 Mappings
 - 1.3 Binary relations
 - 1.4 Orders
 - 1.5 Directed graphs
 - 1.5.1 Directed Acyclic Graphs (DAG's) . .
- 2 Propositional logic
 - 2.1 Subformulae
- 3 Evaluating formulae
 - 3.1 Rewrite rules
 - 3.2 Circuits
- 4 Satisfiability
 - 4.1 Signed Formulae
 - 4.2 Splitting
 - 4.3 Polarity of subformulae
- 5 CNF
- 6 Definitional Clausal Form Translation
- 7 DPLL
- 8 Encoding problems in SAT
- 9 Randomly generated clause sets

10 Randomised algorithms for satisfiability checking

11 Signed Formulae

- 11.1 Finding a model of a specifiable formula . .
- 11.2 Tabelau
- 11.2.1 Branch expansions

1 Setting the stage

Before it is possible to understand some of the content in these notes, it is required to have a basic knowledge of other content that was not covered in my previous notes. This section attempts to give an overview of the relevant topics.

1.1 Decision problems

A decision problem is essentially a question, described by some formal language (such as that of propositional logic) with a yes/no answer. In order to obtain whether the decision problem yields yes or no, input parameters must be specified and the problem evaluated according to the language it is written in.

Another requisite of a decision problem is that the possible set of inputs must be infinite. That is to say that $5 = 5$, or “Is the sky blue?” would not be a decision problem, since the set of inputs in both cases is the \emptyset . Likewise, if a problem doesn’t yield a yes/no answer (for example, a quadratic equation), then it is not a decision problem.

1.1.1 Decidable problems

A problem is classed as decidable if it is both a decision problem and there exists an algorithm that will be able to compute the correct answer for all (infinite number of) inputs to the problem. Such an algorithm can be described

by any Turing complete programming language, in fact, it works both ways; if there is a computer program which can find the correct yes/no answer for a decision problem over all inputs, then the problem must be decidable.

Some decision problems are undecidable. It is impossible to create an algorithm that can always solve the problem for all of its inputs. One such problem is the Halting problem, which asks:

Given the description of an arbitrary program and a finite input, decide whether the program finishes running or will run forever.

It has been proven that the Halting problem is undecidable when running on a Turing machine. The essence of the proof is that the algorithm evaluating whether the input program will halt or not could be made to contradict itself.

1.2 Mappings

A mapping is a function that takes an input from one set and returns an output from another (or the same) set. Conversely, you could describe a function as mapping one set to another.

The notation $f'(x) = f(x) + \{a \rightarrow b\}$ means:

$$f'(x) \stackrel{\text{def}}{=} \begin{cases} b & \text{if } x = a \\ f(x) & \text{otherwise} \end{cases}$$

1.3 Binary relations

Binary relations are often denoted by \Rightarrow , and its reverse is denoted by \Leftarrow such that $y \Leftarrow x \stackrel{\text{def}}{=} x \Rightarrow y$. Along a similar train of thought, the symmetric closure of \Rightarrow is \Leftrightarrow , where $\Leftrightarrow = \Rightarrow \cup \Leftarrow$

1.4 Orders

An order is a relation that is irreflexive and transitive. This means that each pair in the relation cannot be constructed of only one element (e.g. $x > x$ wouldn't be allowed) and if $((x > y) \wedge (y > z)) \implies x > z$.

1.5 Directed graphs

A directed graph consists of a set N and a binary relation on the set R . The elements in N are nodes, and the relation R defines the edges between the nodes. A directed graph is finite if its set is finite.

A *path* is a subset of R where each element of the path will end at the start of another element with the exception of the start and end elements. A cycle is a path where there is no start and end pairs.

1.5.1 Directed Acyclic Graphs (DAG's)

A Directed Acyclic Graph is a directed graph that has no cycles. If a dag has a node n such that every other node in the dag is reachable by n , then the dag is *rooted at n* .

2 Propositional logic

Thankfully, we covered propositional logic in some detail in the COMP11120 course in the first year. However, this module builds upon what was learnt in the first year, and so there is more to learn!

This course introduces the idea of an *interpretation*. An interpretation of a propositional formula is the values that are assigned to each variable in the formula. Assigning values to variables should be a familiar concept, however, calling it an interpretation brings the idea closer to the real world since the formula can be interpreted to mean something in a specific context. Interpretations are also referred to as truth assignments, and boolean values are referred to as truth values.

In this course, we sometimes need to assign a truth value to an expression that contains a conjunction or disjunction and nothing more. If this is the case, then for every interpretation, conjunctions evaluate to true, and disjunctions evaluate to false.

There is a notation to generalise the fact that an interpretation may be a model of many different formulas. If S is

a set of formulas for which an interpretation I is a model, then:

$$I \models S$$

If two formulas can be modelled by exactly the same interpretations, then they are said to be *equivalent*, the symbol for which is \equiv .

For a summary of the above couple of paragraphs, look at the properties in Table 2.

Note:

‘iff’ means ‘if and only if’

Property

A can only be valid if $\neg A$ is unsatisfiable.

A can only be satisfiable iff $\neg A$ is not valid.

A is valid iff $A \equiv \top$.

$A \equiv B$ iff $A \Leftrightarrow B$.

Table 1: Some of the fundamental properties of PL

2.1 Subformulae

A subformula is essentially a formula within a formula. Every formula is composed of one or more subformulae and every formula is it’s own subformula.

Imagine that a formula is composed of layers, with each layer being composed of an operator (such as a number of conjunctions, or an implication). The immediate subformulae of

a formula is the components that the outermost layer uses. It's a hard thing to describe using an informal language such as English, so here are some more concrete examples:

- $A_1 \dots A_n$ are immediate subformulae of $A_1 \vee \dots \vee A_n$, and $A_1 \wedge \dots \wedge A_n$.
- The formula A is an immediate subformula of $\neg A$.
- A_1 and A_2 are immediate subformulae of $A_1 \Rightarrow A_2$ and $A_1 \Leftrightarrow A_2$.
- Every formula is a subformula of itself.
- If A_1 is an immediate subformula of A_2 , and A_2 is an immediate subformula of A_3 , then A_1 is a **subformula** of A_3 .

The above table and previous discussions of equivalence mean that if we replace a subformula A in a formula X by another subformula B that is equivalent to A ($A \equiv B$), to give Y , then $X \equiv Y$.

3 Evaluating formulae

Evaluating formulae can be viewed as a decision problem. If a formula under a specific interpretation is true, then the answer is *yes*, otherwise it's *no*.

It is possible to evaluate formulae relatively efficiently by evaluating each immediate subformula in turn. For example, if we wanted to evaluate the expression $a \Rightarrow b$, we would evaluate b first, since if $b \equiv \top$, then the whole expression

	<i>Subformula</i>				Value
1	$(p \implies q) \wedge (p \wedge q \implies r) \implies$	$(p \implies r)$			1
2			$(p \implies r)$		1
3	$(p \implies q) \wedge (p \wedge q \implies r) \implies$				0
4		$(p \wedge q \implies r)$			1
5	$(p \implies q)$				0
6		$p \wedge q$			0
7	p	p	p		1
8		q	q		0
9			r	r	1

would evaluate to \top . Only in the case where $b \equiv \perp$, would we need to evaluate a .

This method of evaluating formulae can be done cleanly using a table such as the one in Table 3.

3.1 Rewrite rules

It is possible to evaluate a formula simply by rewriting the subformulae inside of it. In order to do this, replace each atom in the formula by it's truth value in the interpretation:

For $A = (x \vee y) \implies z$, and $I = \{x \rightarrow 0, y \rightarrow 1, z \rightarrow 0\}$:

$$\begin{aligned} (x \vee y) &\implies z \\ (\perp \vee \top) &\implies \perp \end{aligned}$$

Then we can use the PL rewrite rules to iterate towards a final truth value:

$$(\perp \vee \top) \implies \perp \equiv \top \implies \perp \quad (1)$$

$$\top \implies \perp \equiv \perp \quad (2)$$

$$\perp \quad (3)$$

It would be very easy to write an algorithm to evaluate formulae using the rewrite rules. Assuming a formula could be represented as a sensible datastructure such as a tree, the algorithm needs to do the following steps:

1. Traverse the formula and replace all atoms by their values in the current interpretation.
2. While the formula is not \top or \perp , re-write the formula using the re-write rules
3. Return either \top or \perp

3.2 Circuits

There exists a certain isomorphism between PL and digital circuitry. It is often possible to reduce a problem of digital logic (one that involves the use of AND, OR, XOR and other gates) into one that can be represented by a PL formula.

Though this may seem both trivial and frivolous at first, there exist many good reasons for such an exercise. For example, it would be possible to check if two circuits did the same thing using techniques from PL, or you could use

re-write rules to make one circuit into another equivalent, yet more compact circuit.

4 Satisfiability

Checking if a PL formula is satisfiable (there is at least one interpretation where the formula is true) is a famous decision problem.

The satisfiability problem is NP-Complete (so we can check the solution quickly, but we cannot arrive at the solution in polynomial time).

One simple way of checking the satisfiability of a formula is to enumerate all the possible interpretations, and put them in a truth table. As usual, you split the formula into its subformulae, but unlike a conventional truth table, you have a column for each interpretation (so enumerate the possible combinations of all the atomic formulae).

Using truth tables as a method of satisfiability checking is not very efficient. Because there is a requirement to enumerate all of the interpretations, you need to do as many operations as there are $2^{\#atoms}$.

When you think about it, there will be certain situations, where we don't have to enumerate all of the different interpretations. For instance if we assume that $I(r) = \top$, then we know that $I(p \vee r) = \top$ and $I(x \implies (p \vee r)) = \top$. In this example, we could eliminate lots of possible interpretations since we know some of them to be valid without even trying them.

If we were to take this principle into account when evaluating formulae for satisfiability, how many enumerations we would need to do would depend on which variables we enumerated first. Good SAT checkers will choose wisely in this regard.

4.1 Signed Formulae

In order to progress, we need to understand the notion of signed formulae.

A signed formula, $A = b$ is one where the formula A is equal to b in a specific interpretation I . It follows that:

- For any A and any I either $A = 1$ or $A = 0$ in I .
- The formula A is satisfiable iff $A = 1$

4.2 Splitting

Splitting is another way of checking satisfiability of formulae. The idea is that we consider truth values for variables in the formula, and simplify the formula based on those considerations. Formally:

For every propositional formula A , and every atom p , denote A_p^\top and A_p^\perp as the formulas with p replaced with \top and \perp respectively.

The splitting algorithm tells us if each interpretation of a formula is satisfiable or unsatisfiable. If all the interpreta-

tions are \top , then the formula must be valid, if they are all \perp , then it must be unsatisfiable.

Figure 1, (the same example used in Andrei's notes) shows the splitting algorithm in action. I think its easiest to learn splitting through diagrams such as this, and thus I won't write any more about how it works.




Figure 1: Splitting the formula $\neg((p \implies q) \wedge (p \wedge q \implies r) \implies (p \implies r))$

4.3 Polarity of subformulae

5 CNF

6 Definitional Clausal Form Translation

7 DPLL

8 Encoding problems in SAT

9 Randomly generated clause sets

10 Randomised algorithms for satisfiability checking

11 Signed Formulae

A signed formula is one where there is an expression, A and a boolean b , where $A = b$. If $A = b$ is *true*, in an interpretation I , then it is denoted by $I \models A = b$, and consequently, I is a model of the signed formula $A = b$.

Note:

This is also when $I(A) = b$.

If a signed formula has a model, then it is specifiable.

11.1 Finding a model of a specifiable formula

If we had a signed formula such as $A \Leftrightarrow B = 1$, the three interpretations that model it are:

A	B
0	0
0	1
1	1

11.2 Tabelau

A tableau is a tree with each node being a signed formula. The tableau for the signed formula $A = b$ would have the root node as $A = b$.

The notation for a set of branches is $B_1 | \dots | B_n$, where each B_i is a branch.

11.2.1 Branch expansions

There are a number of rules that can be used to expand the branches of a tableau.