

# COMP28112 Model Answers

Todd Davies

May 20, 2015

## 2011 Paper

**Describe one key difference between client-server and peer-to-peer applications.** 2011.1.a  
(2 marks)

P2p does not have a central point of coordination; all nodes are the same in terms of functionality (though they may have different roles). A client server architecture is dissimilar to this in the fact that the machines that make up the network have different capabilities and responsibilities that cannot be changed.

**Explain briefly what failures are known as Byzantine failures and what the Byzantine generals problem refers to.** 2011.1.b  
(2 marks)

Byzantine failures are ones where processes send failed messages since they are malicious or faulty.

The Byzantine Generals problem refers to a situation where two generals are situated on the sides of a valley and want to attack a city in the valley. To win the battle, they must attack at the same time, but they cannot coordinate their attack with any guarantee of timing it properly since they can only send messages through the city where the messengers may be killed or the contents of the message changed.

**Describe briefly the two-phase commit protocol.** 2011.1.c  
(2 marks)

**Explain briefly why somebody would choose to use Cloud Computing.** 2011.1.d  
(2 marks)

**Why is it practically impossible to achieve strict consistency in a distributed system?** 2011.1.e  
(2 marks)

It is practically impossible to achieve strict consistency in a distributed system since latency is not zero, so you can never know exactly what is going on in a remote system at any one time. Messages can get lost or mutated too, which further complicates the sending and updating of state between systems.

**Suppose the C function below is to be made available to remote processes using RPC. What particular implementation problem does this highlight? Why is the normal solution only partly satisfactory?** 2011.1.f  
(2 marks)

```
void doIt (int *p, int *q) { (*p)++ ; (*q)-- ; return ; }
```

You would have to serialise and deserialise the values of p and q to have this function run over RPC, because it uses pointers as opposed to actual values. You would also have to update the value of the pointers either as they are updated during the RPC functions execution or after the function has finished executing.

**Distinguish carefully between a name server and a directory server.**

2011.1.g  
(2 marks)

A name server takes the name of a resource and resolves it into the address of the resource (e.g. a URL). A directory server takes attributes about an object (e.g. the amount of RAM of a server or the age and course studied of a student) and returns information about objects that have the same attributes (think LDAP)

**What is meant by the term causally ordered multicast?**

2011.1.h  
(2 marks)

Ummmmm? Maybe take a look here: <http://www.cse.buffalo.edu/~stevko/courses/cse486/spring13/lectures/12-multicast2.pdf>

Looks like every process has a vector clock and when a message is received, the process waits until it can preserve the causal ordering before accepting the message. This means it waits until messages that had been previously sent to it (according to the vector clock) arrive before accepting the new message. Until then, the message is buffered.

**What properties are provided by a secure channel?**

2011.1.i  
(2 marks)

A secure channel is one that third parties cannot eavesdrop on. They might be able to see the data passing through, but it will be encrypted somehow so they won't be able to make sense of it.

**The Andrew File System (AFS) uses callback promises. Explain what this means.**

2011.1.j  
(2 marks)

Dunno? Google turns up not much...

Im guessing that it is a mechanism where if you send a remote server a message, it promises to reply back to you and will cause a callback to execute in your code when it does.

**Explain briefly why some applications are not parallelisable. Describe Amdahls law and explain what it can be used for.** 2011.2.a  
(3 marks)

Some applications are not parallelisable because they have operations that are interdependent. For example, if we were building a machine to compute how long it takes for the collatz conjecture to reach 1 given an input number, we could not parallelise it since each stage of the conjuncture depends on the last. Amdahls law dictates how much we can speed up the execution of a program given how much of the program is parallelisable. The law is:

$$\frac{1}{\text{serial portion} + \left(\frac{1}{\text{num threads}} \times \text{parallel portion}\right)}$$

**Explain briefly what the four properties commonly denoted by the acronym ACID are when referring to transactions.** 2011.2.b  
(3 marks)

ACID stands for:

*Atomic* - All or nothing principle; either the transaction is committed and its changes are applied, or no state is changed as a result of the transaction.

*Consistent* - Each transaction moves the system from one valid and consistent state to another.

*Isolated* - Each transaction executes independently from all other transactions (even if they may be happening concurrently). The effect of transactions in progress is hidden from all other transactions.

*Durable* - Once a transaction is committed, it stays committed even in the event of failure such as power loss etc

**A service is replicated onto 3 computers.** 2011.2.c  
(6 marks)

- The first computer, A, has a mean time between failures of 2 days.
- The second computer, B, has a mean time between failures of 3.5 days.
- The third computer, C, has a mean time between failures of 12 days.

**When a failure occurs, it takes on average 12 hours to fix.**

**What is the availability of the replicated service?** 2011.2.c.i  
(2 marks)

We need to find the chance that all the computers will fail at the same time. It is easy to find the up-times for each individual computer:

A - 75%

B - 85.7%

C - 95.83%

The chance that these will all fail at the same time is:  $(1 - 0.75) * (1 - 0.875) * (1 - 0.9853) = 0.00046 = 0.046\%$   
Therefore, the uptime is going to be  $100 - 0.046 = 99.954\%$

**What would the availability of the replicated service be if only computers A and B were used?** 2011.2.c.ii  
(2 marks)

If only A and B were used, it would be:

$$1 - ((1 - 0.75) * (1 - 0.875)) = 0.96875 = 96.88\%$$

**Describe how in the general case of n computers, each with a mean time between failures  $f_i$  and an average time to fix a failure  $t_i$  you would choose the two computers that provide the highest availability.** 2011.2.c.iii  
(2 marks)

Rank the computers in ascending order of  $t_x/f_x$ , and take the top  $n$  elements.

The following four processes access a shared variable  $x$ . Each process accesses a different replica of the store used to hold this variable. Before any process starts executing, the value of  $x$  is 0 in all the replicas. 2011.2.d  
(8 marks)

<u>Process 1</u>	<u>Process 2</u>	<u>Process 3</u>	<u>Process 4</u>
$x=1;$	$x=2;$	$while(x==0);$	$while(x==0);$
$x=4;$	$x=3;$	$y=x;$	$z=x;$
		$y=5*x+y;$	$z=5*x+z;$

When all four processes have completed executing the statements given, are 7 and 14 possible values of  $y$  and  $z$  respectively, if the replication uses the sequential consistency model? Justify your answer. 2011.2.d.i  
(4 marks)

With a sequential consistency model, it is not possible for the final values to be  $y = 7$  and  $z = 14$ , since for  $y$  to equal 7,  $x$  must change state while process 3 is running, and this is not possible in a sequential consistency model since only one process can execute at once.

The possible values of  $x$  when process three actually starts are  $x = 4$  and  $x = 3$ , none of which let  $y = 7$ .

When all four processes have completed executing the statements given, are 7 and 14 possible values of  $y$  and  $z$  respectively, if the replication uses the causal consistency model? Justify your answer. 2011.2.d.ii  
(4 marks)

Using a causal consistency model, we still can't get  $y = 7, z = 14$ , since we need  $y = 5 \times (x = 1) + (y = 2)$  and  $z = 5 \times (x = 2) + (z = 4)$ , but this can never happen, since for  $z = 4$ , we need to have done  $x = 1$ , but we need to have done  $y = 2$  before that, which means that we would have already done  $x = 2$ .

**Describe in detail the Bully algorithm for the election of a leader in a distributed system.** 2011.3.a.i  
(8 marks)

The Bully Algorithm works like this:

- An initiating client will send a message to all other clients with a higher identifier than itself.
- Any client receiving an election message will then start its own election, sending messages to clients with higher identifiers than itself.
- Only when a process gets no replies (within a timeout) is it considered elected. It then sends a message to all other processes announcing its election.

In this manner, any process that receives a message should always send a message back to the sender

**Carefully explain all the assumptions made in the Bully algorithm.** 2011.3.a.ii  
(4 marks)

- The timeouts are reasonable
- All processes know about all other processes (and their loads)
- Messages are delivered quickly
- Messages are delivered reliably

**Describe one application for which the Bully algorithm might be applied, indicating why a leader is needed.** 2011.3.b  
(2 marks)

In a peer-to-peer protocol (such as bit torrent), the Bully algorithm can be used to elevate a node to coordinator status if the client that was coordinating the network went offline or started to get a too high load.

**In a system containing 6 computers, identified by the integers 1-6, the leader is chosen by the Bully algorithm to be the live one with the highest identifier. Assume for this part that all messages are delivered promptly, and that the computers and the network are entirely reliable.** 2011.3.c  
(6 marks)

**How many messages in total are sent so that the computer with identifier 1 after it is rebooted can learn the identity of the leader by triggering an election? Take care to explain your working!** 2011.3.c.i  
(6 marks)

Since identifier 1 is the lowest node, it will exhibit worst case behaviour for the bully algorithm,  $O(n^2)$ .

This is because 1 will send five messages out (to each other process) to start the election. Client 2 will then send out four messages (to the ones bigger than it) etc, until there are  $5 + 4 + 3 + 2 + 1 + 0 = 15$  messages sent. Since every client receiving an election message replies to the sender, then 15 replies will be sent (in addition to the 15 initial message). Then, the winner of the election (6) will send a message out to say that it is elected, which is another 5 messages, bringing the total to 35.

2011.3.c.ii  
(6 marks)

If 5 wanted to find the leader, it would send a message to 6 and get a reply that 6 is online, and another one saying 6 is the coordinator, which is three messages.

## 2012 Paper

**What is a Java servlet?**

2012.1.a  
(2 marks)

A Java Servlet is a Java program with the capabilities of a server. It could host web pages, provide an API endpoint or other services and is most commonly used to serve content via the HTTP protocol.

**What is the main assumption on which Cristian's clock synchronisation algorithm is based?**

2012.1.b  
(2 marks)

The time for a message to go from machine A to machine B is (roughly) the same as the time it would take for a message to go from machine B to machine A. This is most often accurate for routes with small round trip times.

**Explain the difference between a name server and a directory server.**

2012.1.c  
(2 marks)

A name server takes a name, matches it to an object and returns attributes about the object. A directory server takes attributes, matches them to an object and returns more attributes about the object.

**Explain briefly what is meant by the term middleware.**

2012.1.d  
(2 marks)

Middleware is software that sits between a client application and the operating system. It provides services to the client application that the operating system does not, such as RPC stubs.

It can also provide an abstraction from the OS to mask the heterogeneity of platforms used in distributed systems (some apps will run on Windows, others on Linux, some on x64, some on ARM architectures etc).

**Why is it practically impossible to achieve strict consistency in a distributed system?**

2012.1.e  
(2 marks)

Strict consistency is when any read to a shared data item returns the most recent write operation on that data item. This means there must be an absolute time ordering of all accesses. Unfortunately, since (as we know from the eight fallacies of Distributed Computing), the latency of any network is not zero and messages are not reliable. This means that any message we send to update other machines about a change of state may not be sent. Since there is no way of getting an absolute global clock or getting any global state, then we cannot achieve real time memory consistency across all nodes, and therefore cannot achieve strict consistency in the system.

**Traditional RPC mechanisms cannot handle pointers. What is the problem?**

2012.1.f  
(2 marks)

In order to handle pointers with RPC, you must serialise the datastructure into a message so that it can be constructed from the same message at the other machine. When the reply is received back at the sender, it can be deserialised again and the values in the datastructure that were changed on the remote machine can be updated in local memory. Traditional RPC mechanisms didn't have this facility.